

# Key-Value Stores

Big Data Management

# Definition

A **key-value** store is a simple NoSQL database using an associative array (similar to a map or dictionary) as its core data model. In this model, each key maps to exactly one arbitrary value within a collection.


- No query language
- Simple operations: *get*, *put*, and *delete*

**Wide-column** stores are key-value stores that further structure the value into families, which contain groups of attributes (a.k.a. columns).

# Definition


- Key-value stores

- Entries in form of key-values
  - One key maps only to one value
- Query on key only



- Wide-column stores (Bigtable)

- Entries in form of key-values
  - But now values are split in wide-columns
- Typically query on key
  - May have some support for values



# Bigtable: A Distributed Storage System for Structured Data

FAY CHANG, JEFFREY DEAN, SANJAY GHEMAWAT, WILSON C. HSIEH,  
DEBORAH A. WALLACH, MIKE BURROWS, TUSHAR CHANDRA,  
ANDREW FIKES, and ROBERT E. GRUBER  
Google, Inc.

---

Bigtable is a distributed storage system for managing structured data that is designed to scale to a very large size: petabytes of data across thousands of commodity servers. Many projects at Google store data in Bigtable, including web indexing, Google Earth, and Google Finance. These applications place very different demands on Bigtable, both in terms of data size (from URLs to web pages to satellite imagery) and latency requirements (from backend bulk processing to real-time data serving). Despite these varied demands, Bigtable has successfully provided a flexible, high-performance solution for all of these Google products. In this article, we describe the simple data model provided by Bigtable, which gives clients dynamic control over data layout and format, and we describe the design and implementation of Bigtable.

Categories and Subject Descriptors: C.2.4 [**Computer Communication Networks**]: Distributed Systems—*distributed databases*

General Terms: Design

Additional Key Words and Phrases: Large-Scale Distributed Storage

## Bigtable: A Distributed Storage System for Structured Data

Use cases where:

We want to scale to a very large size: petabytes of data across thousands of commodity servers.

Google Use Cases:

- **Web Indexing:**
  - Data: URLs, pages ...
  - Latency requirement: backend bulk processing
- **Google Earth:**
  - Data: Satellite images
  - Latency requirement: real time data serving

...

# BigTable Data Model

"A BigTable is a **sparse, distributed, persistent, multi-dimensional, sorted map**."

F. Chang et al.

- Sparse: most elements are unknown
- Distributed: cluster parallelism
- Persistent: disk storage (GFS)
- Multi-dimensional: values with columns
- Sorted: sorting lexicographically by primary key
- Map: lookup by primary key (a.k.a. finite map)

# map

- At its core BigTable is a **map**
  - In other programming languages: associative array (PHP), dictionary (Python), Hash (Ruby), or Object (Javascript)

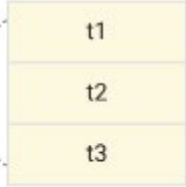
map

```
{  
  "zzzzz" : "woot",  
  "xyz" : "hello",  
  "aaaab" : "world",  
  "1" : "x",  
  "aaaaa" : "y"  
}
```

# sparse

- If a column is not used in a particular row, it does not take up any space. Optional values are not represented by placeholders or input parameters. Memory efficient.

	Column family 1		Column family 2	
	Column 1	Column 2	Column 1	Column 2
Row key 1				
Row key 2				





# distributed

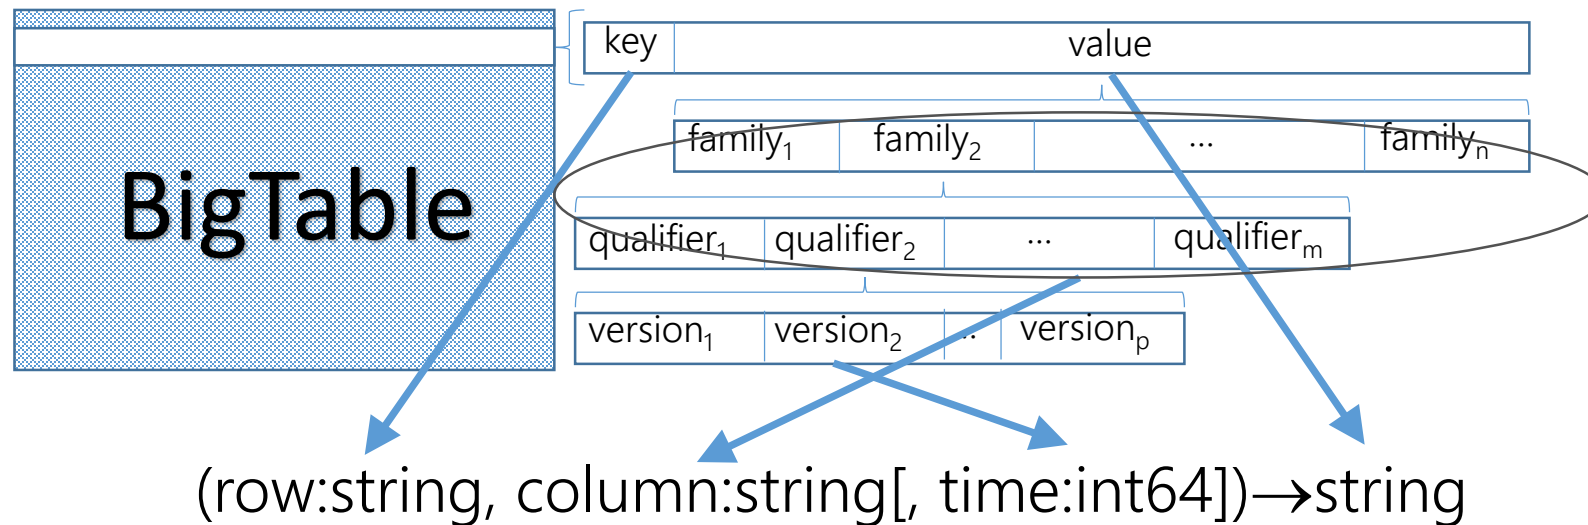
- BigTable are built upon distributed file systems, so that the underlying file storage can be spread out among an array of independent machines
  - BigTable makes use of the GFS

# persistent

- Persistent means that the data you put in the special map “persists” after the program that created or accessed it finished
  - E.g., Similar to a file in a file system

# multidimensional

- Different 'levels' in the map
- Each cell can contain multiple versions of same data
  - (Row Key) x (Column Key) x (Timestamp) → Value
  - Column Key → family:qualifier



# sorted







- Bigtable maintains data in lexicographic order by row key which are indexed
- The row range for a table is dynamically partitioned; each called a **tablet**
- A tablet is the unit of distribution and load balancing in BigTable

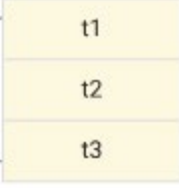
## sorted map

```
{  
  "1" : "x",  
  "aaaaa" : "y",  
  "aaaab" : "world",  
  "xyz" : "hello",  
  "zzzzz" : "woot"  
}
```

# BigTable Data Model

- Stores data in tables composed of indexed rows identified by **row keys**
  - Each row is a set of key-value pairs, where each key is defined by the combination of **column family**, **column qualifier**, and **timestamp**.
- Data is stored as uninterpreted strings (no predefined data types)
- Each cell can store multiple versions of data, each version identified by a unique **timestamp**.

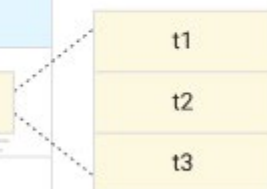
	Column family 1		Column family 2	
	Column 1	Column 2	Column 1	Column 2
Row key 1				
Row key 2				



# BigTable Data Model

- Columns families contain one or more often related columns.
- Families are explicitly created when table is created or updated.
- Qualifiers are column names within the column family.
- Qualifiers are declared at insertion time.

	Column family 1		Column family 2	
	Column 1	Column 2	Column 1	Column 2
Row key 1				
Row key 2				



# HBase

- Apache project
  - Based on Google's Bigtable
- Designed to meet the following requirements
  - Access specific data out of petabytes of data
  - It must support
    - Key search
    - Range search
    - High throughput file scans
  - It must support single row transactions (mutations, such as read, write, and delete requests, are always atomic at the row level)

# BigTable - Hbase Comparison

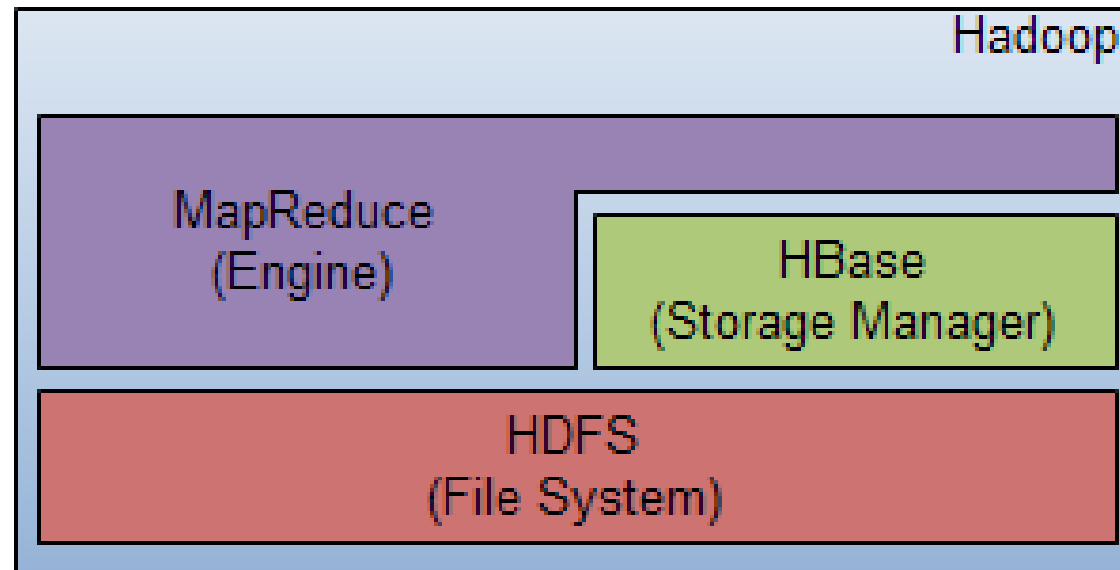
Feature	BigTable	HBase
Origin	Proprietary (Google)	Open-source (Apache)
Infrastructure	Cloud based	Self managed
Distribution unit	Tablets	Regions
Storage Backend	GFS (Colossus)	HDFS
Operational cost	Low (fully managed)	High (manual)



# HBase Architecture

# Hadoop Ecosystem

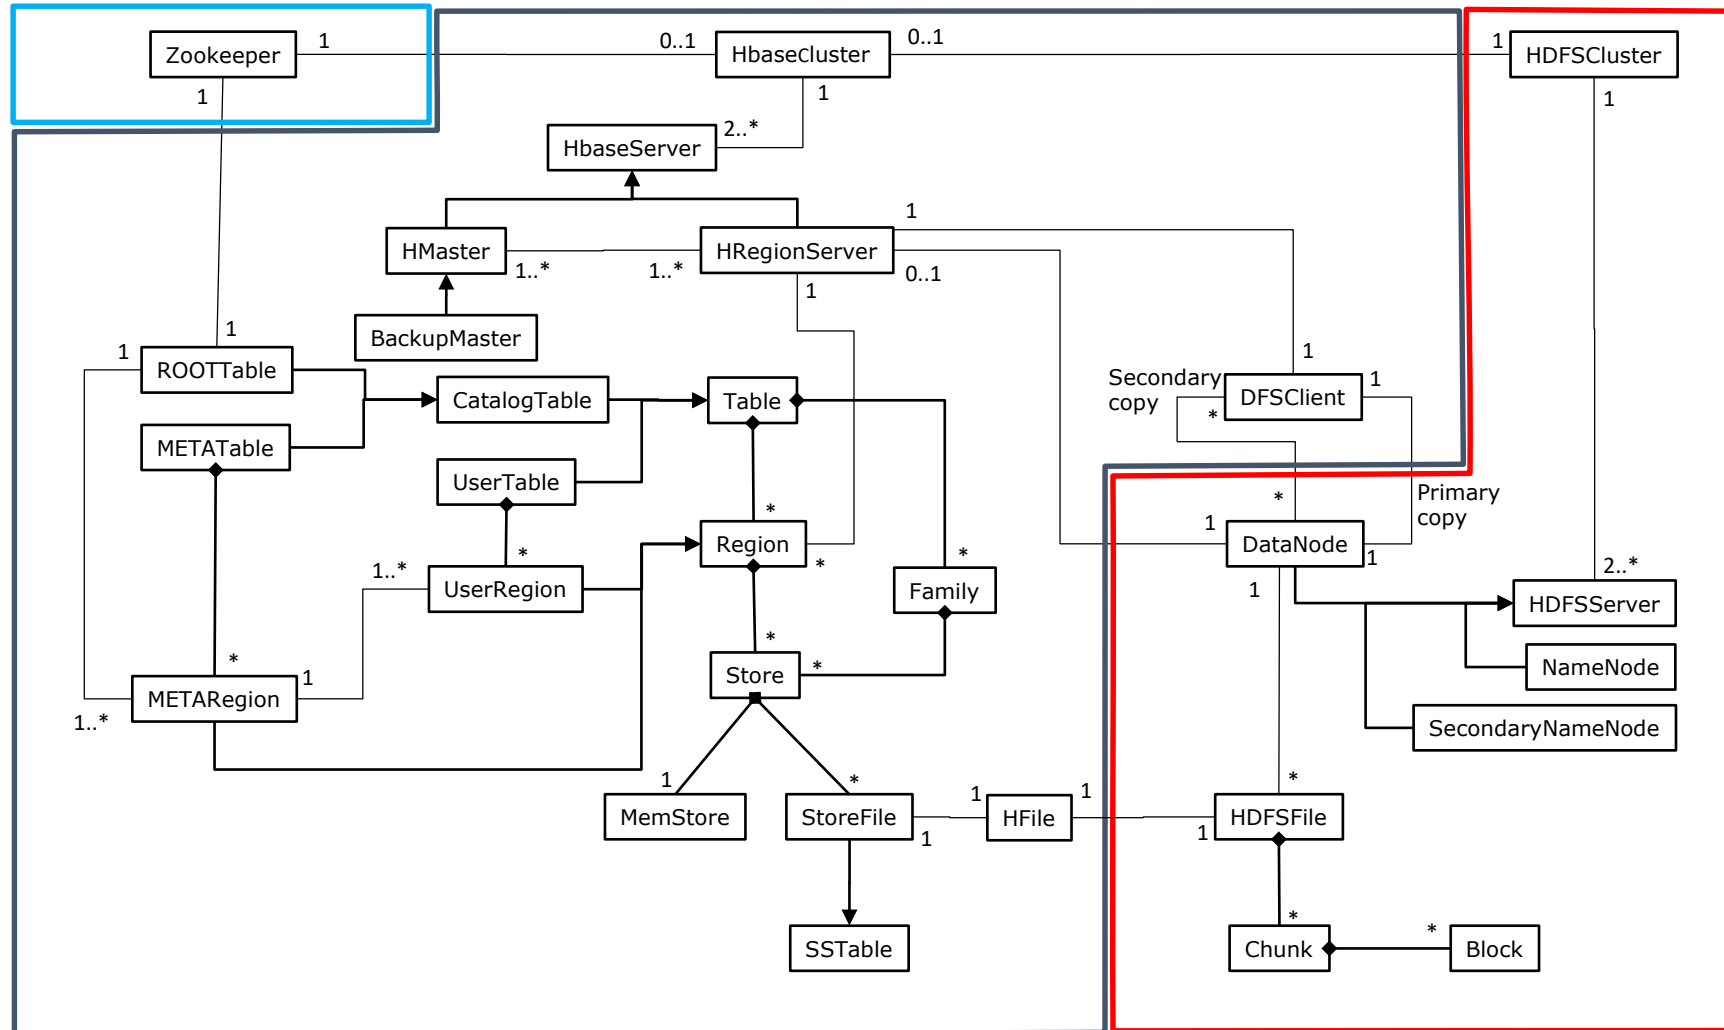
- HDFS physically stores the data in files
- HBase manages the data, and provides basic indexing capabilities allowing random access
- MapReduce is the query execution engine that filters, transforms and aggregates these data.



# HBase shell

- ALTER <tablename>, <columnfamilyparam>
- COUNT <tablename>
- CREATE TABLE <tablename>
- DESCRIBE <tablename>
- DELETE <tablename>, <rowkey>[, <columns>]
- DISABLE <tablename>
- DROP <tablename>
- ENABLE <tablename>
- EXIT
- EXISTS <tablename>
- GET <tablename>, <rowkey>[, <columns>]
- LIST
- PUT <tablename>, <rowkey>, <columnid>, <value>[, <timestamp>]
- SCAN <tablename>[, <columns>]
- STATUS [{summary|simple|detailed}]
- SHUTDOWN

# Functional components

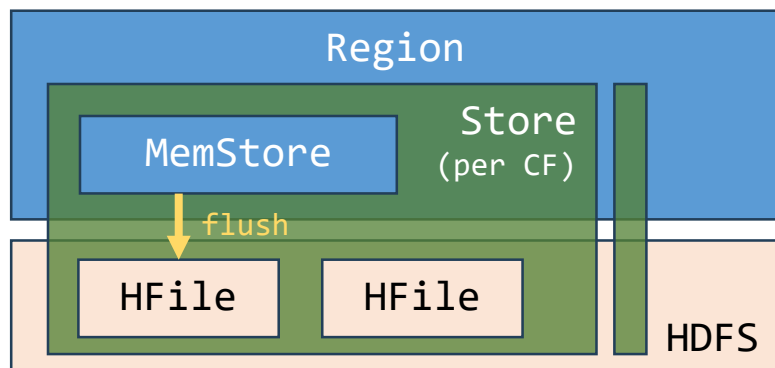


# Functional components

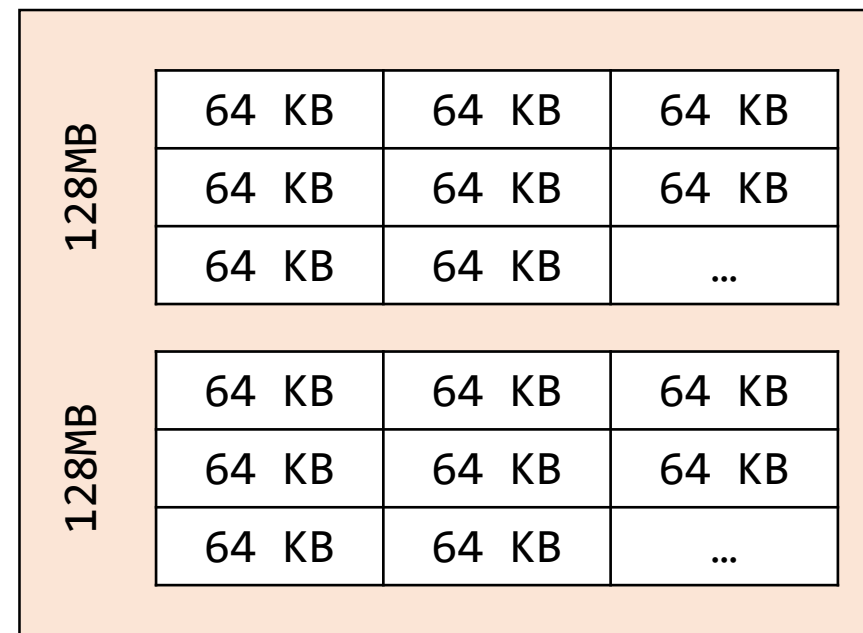
- **Zookeeper** - group of servers that stores HBase system config info
- **HMaster** - Coordinates splitting of regions/rows across nodes
  - Controls distribution of HFile chunks
- **HRegionServer** - Region Servers
  - Serves HBase client requests
    - Manage stores containing all column families of the region
  - Logs changes
  - Guarantees “atomic” updates to one column family

# StoreFiles

- When the MemStore is full (128MB), data are flushed to HDFS
- A StoreFile is generated
  - Format **HFile**
- An HFile stores data into HDFS chunks
  - Chunks are structured into HBase blocks
    - Size 64 KB



Storefile  
(HFile format)



# Functional components

- Tables → Regions → Stores
- Regions are the minimal distribution unit used by HBase
- Data are physically stored in stores:
  - in in-memory buffers (memstores)
  - flushed to disk as storefiles (a.k.a. SStables)
  - storefiles are represented as HFiles (with added metadata)
- HFiles are divided into HBase blocks (64KB)
- HFiles are then written to HDFS and replicated

# HBase processes

## a) Flush

- On memory structure reaching threshold
- Takes memory content and store it into an SSTable (HFile)
- Generates different disk versions of the same record

## b) Minor compactation

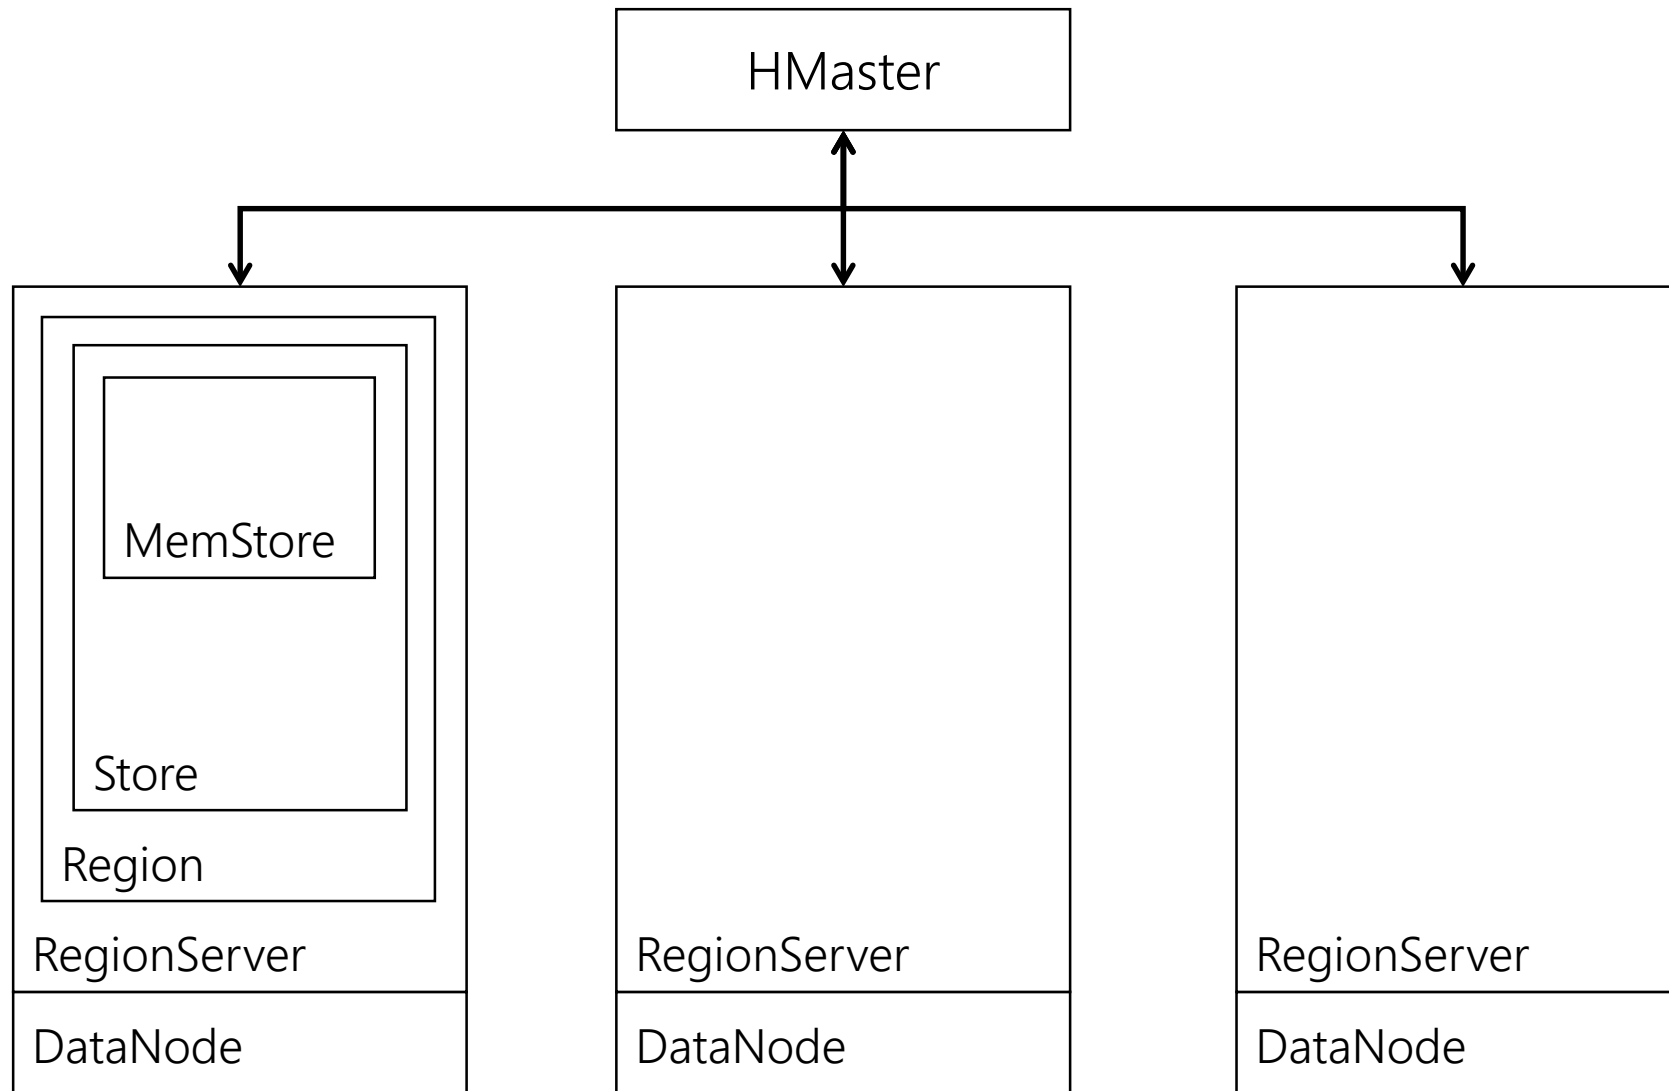
- Runs regularly in the background
- Merges a given number of equal size SSTables into one
- Does not remove all record versions (only some)

## c) Major compactation

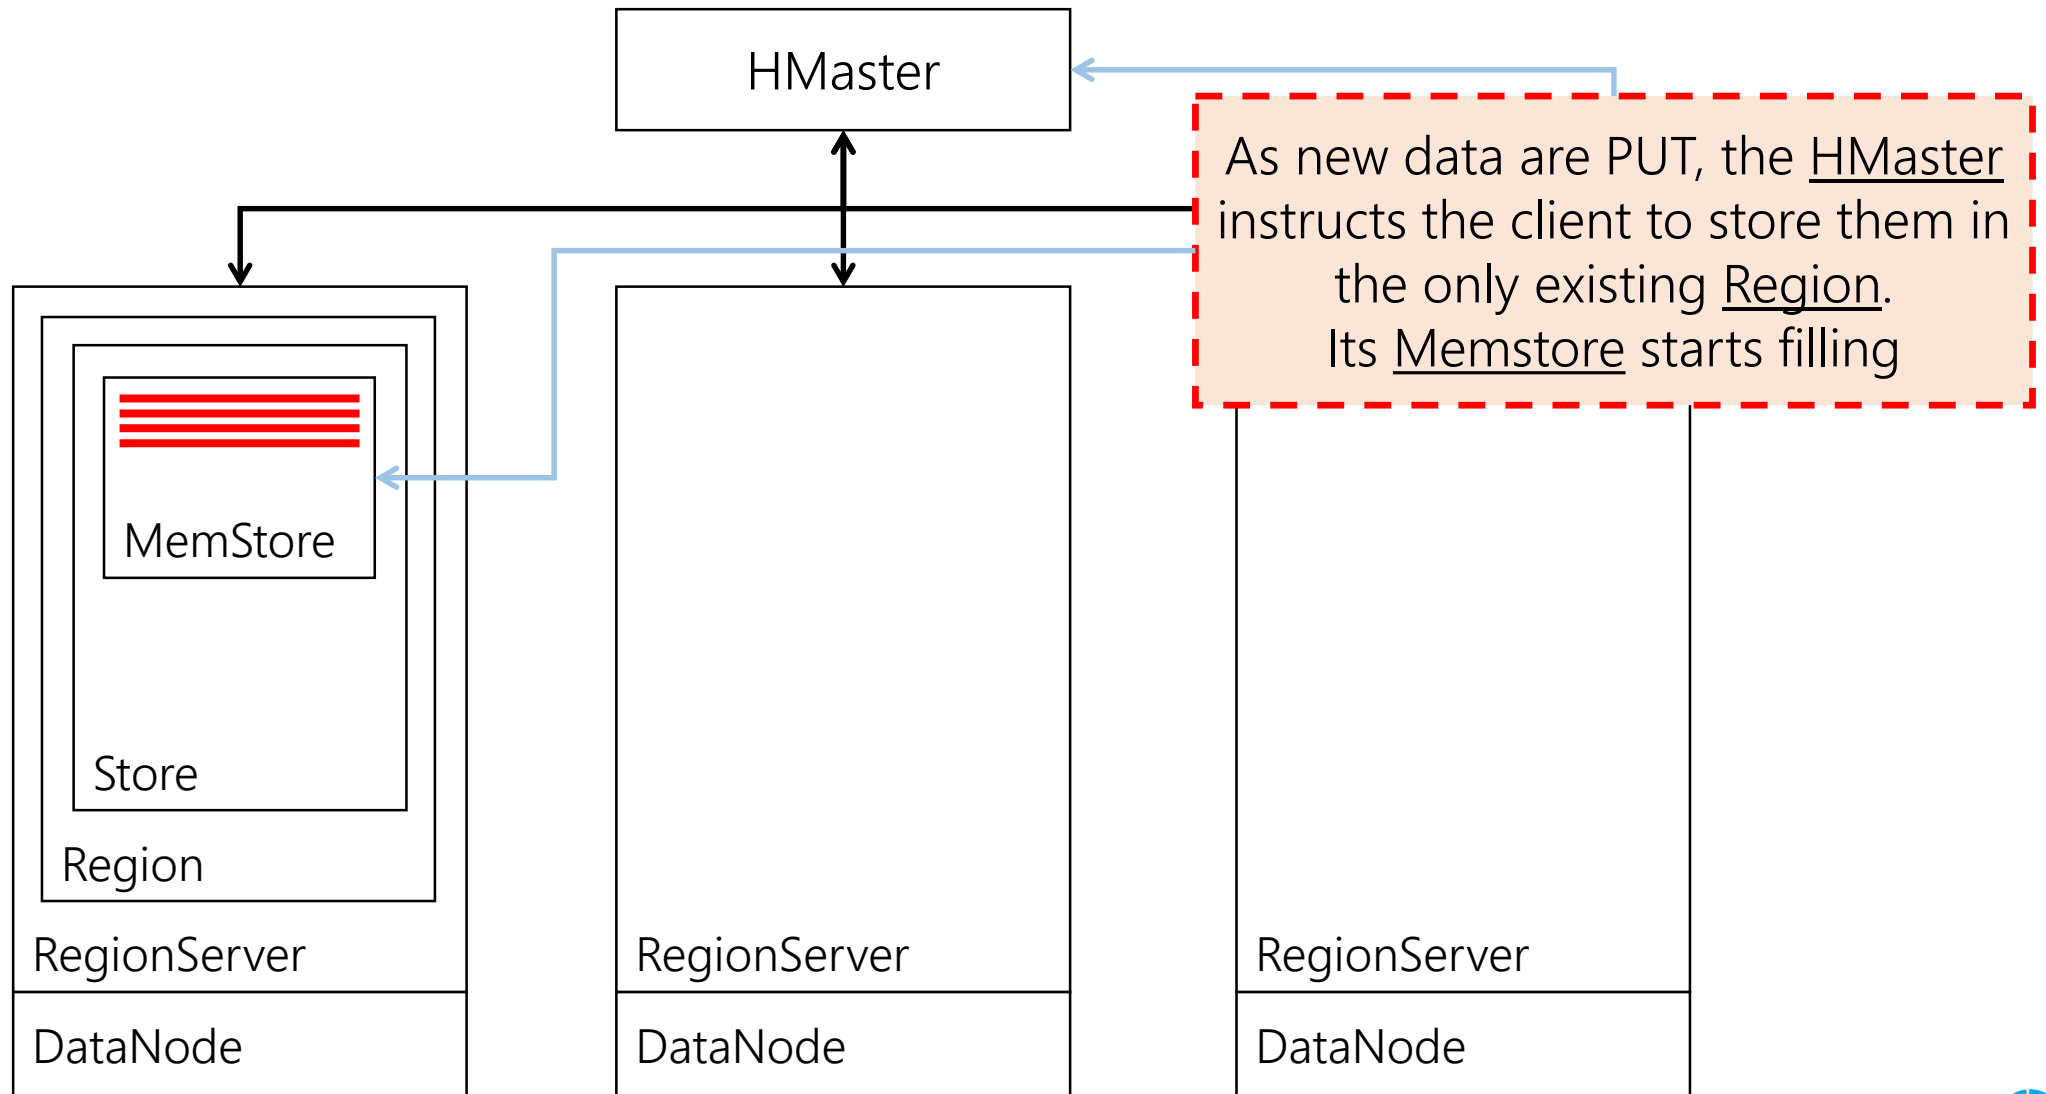
- Triggered manually
- Merges all SSTables
- Leaves one single SSTable
- All versions of a record are merged into one



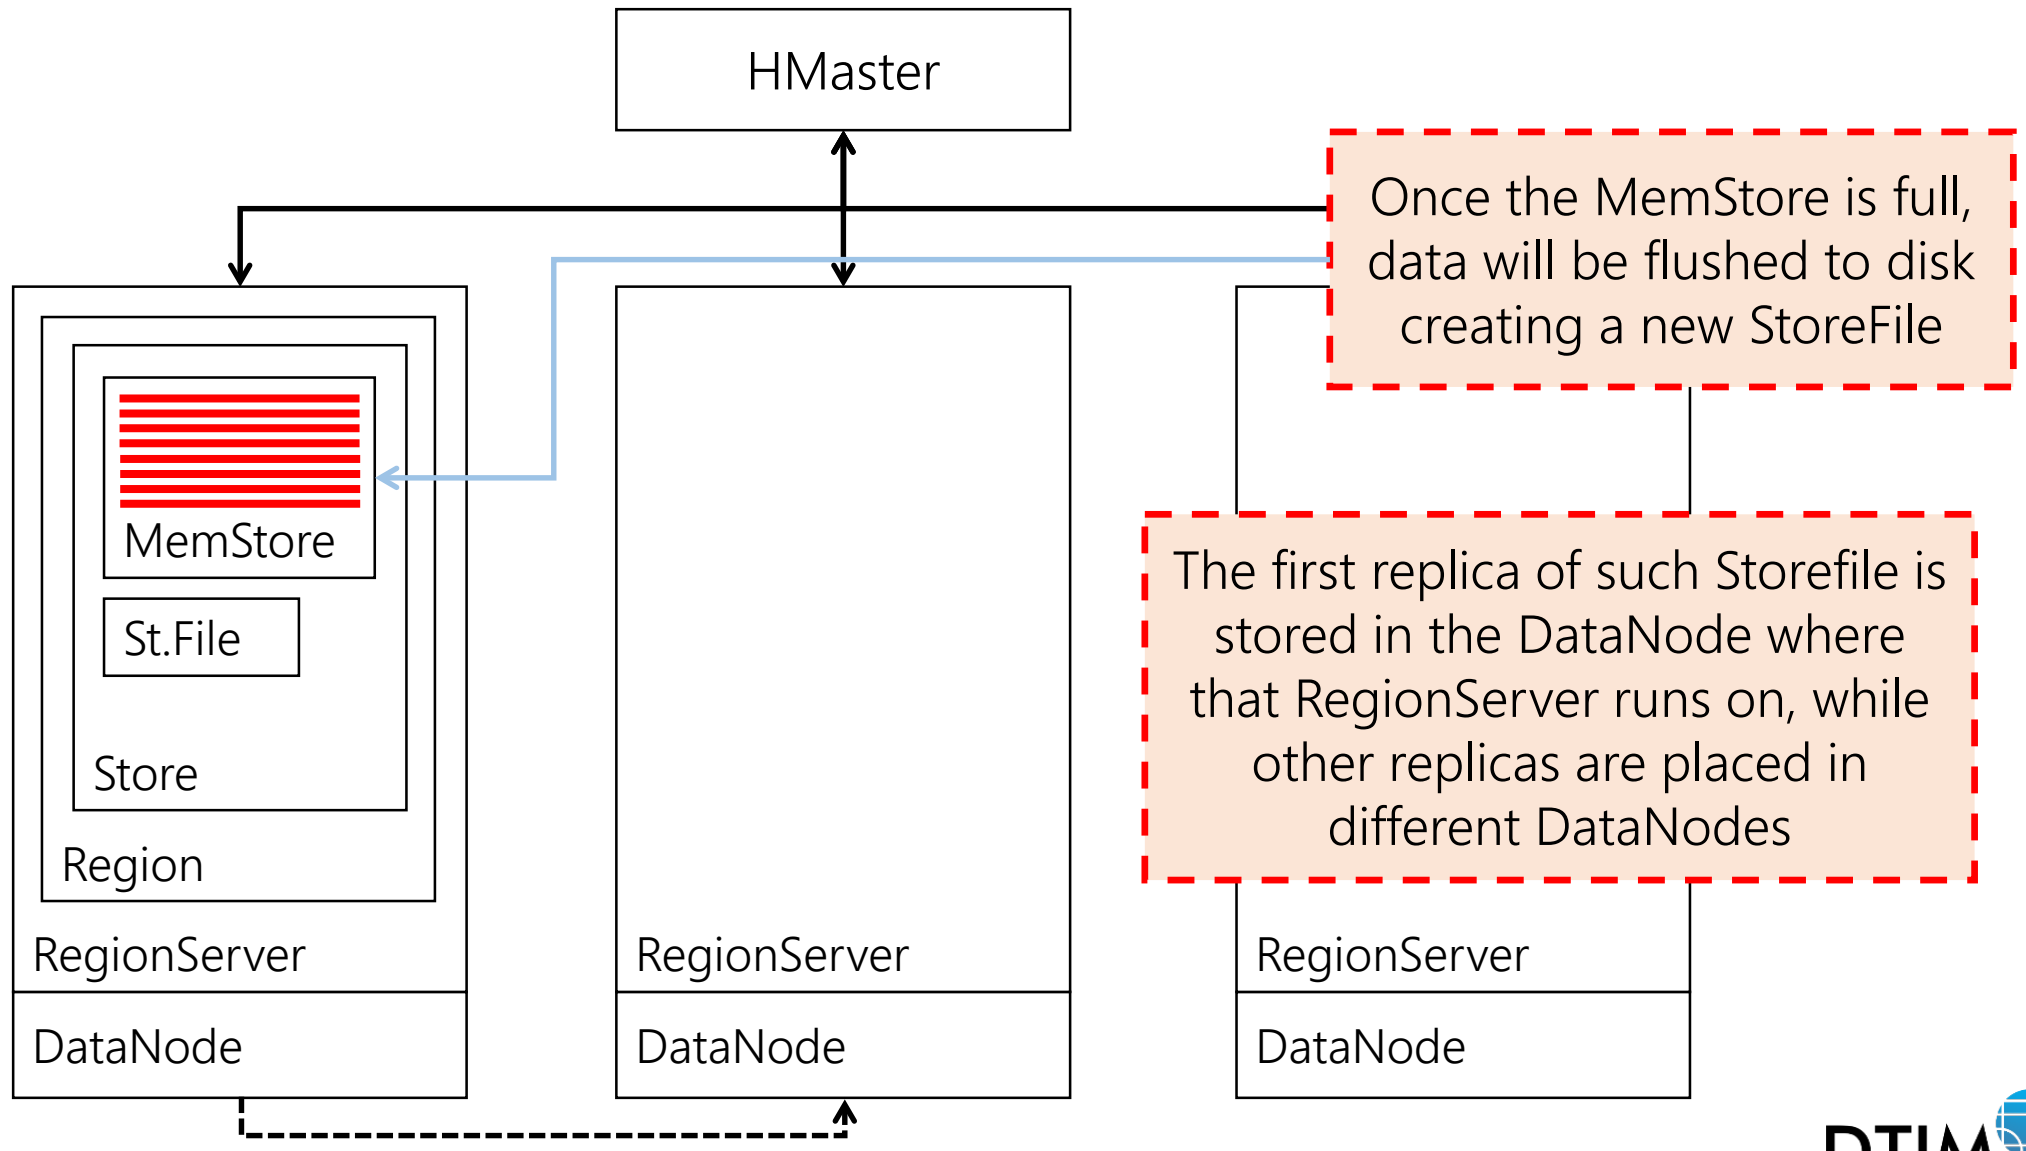
# Example of Flush



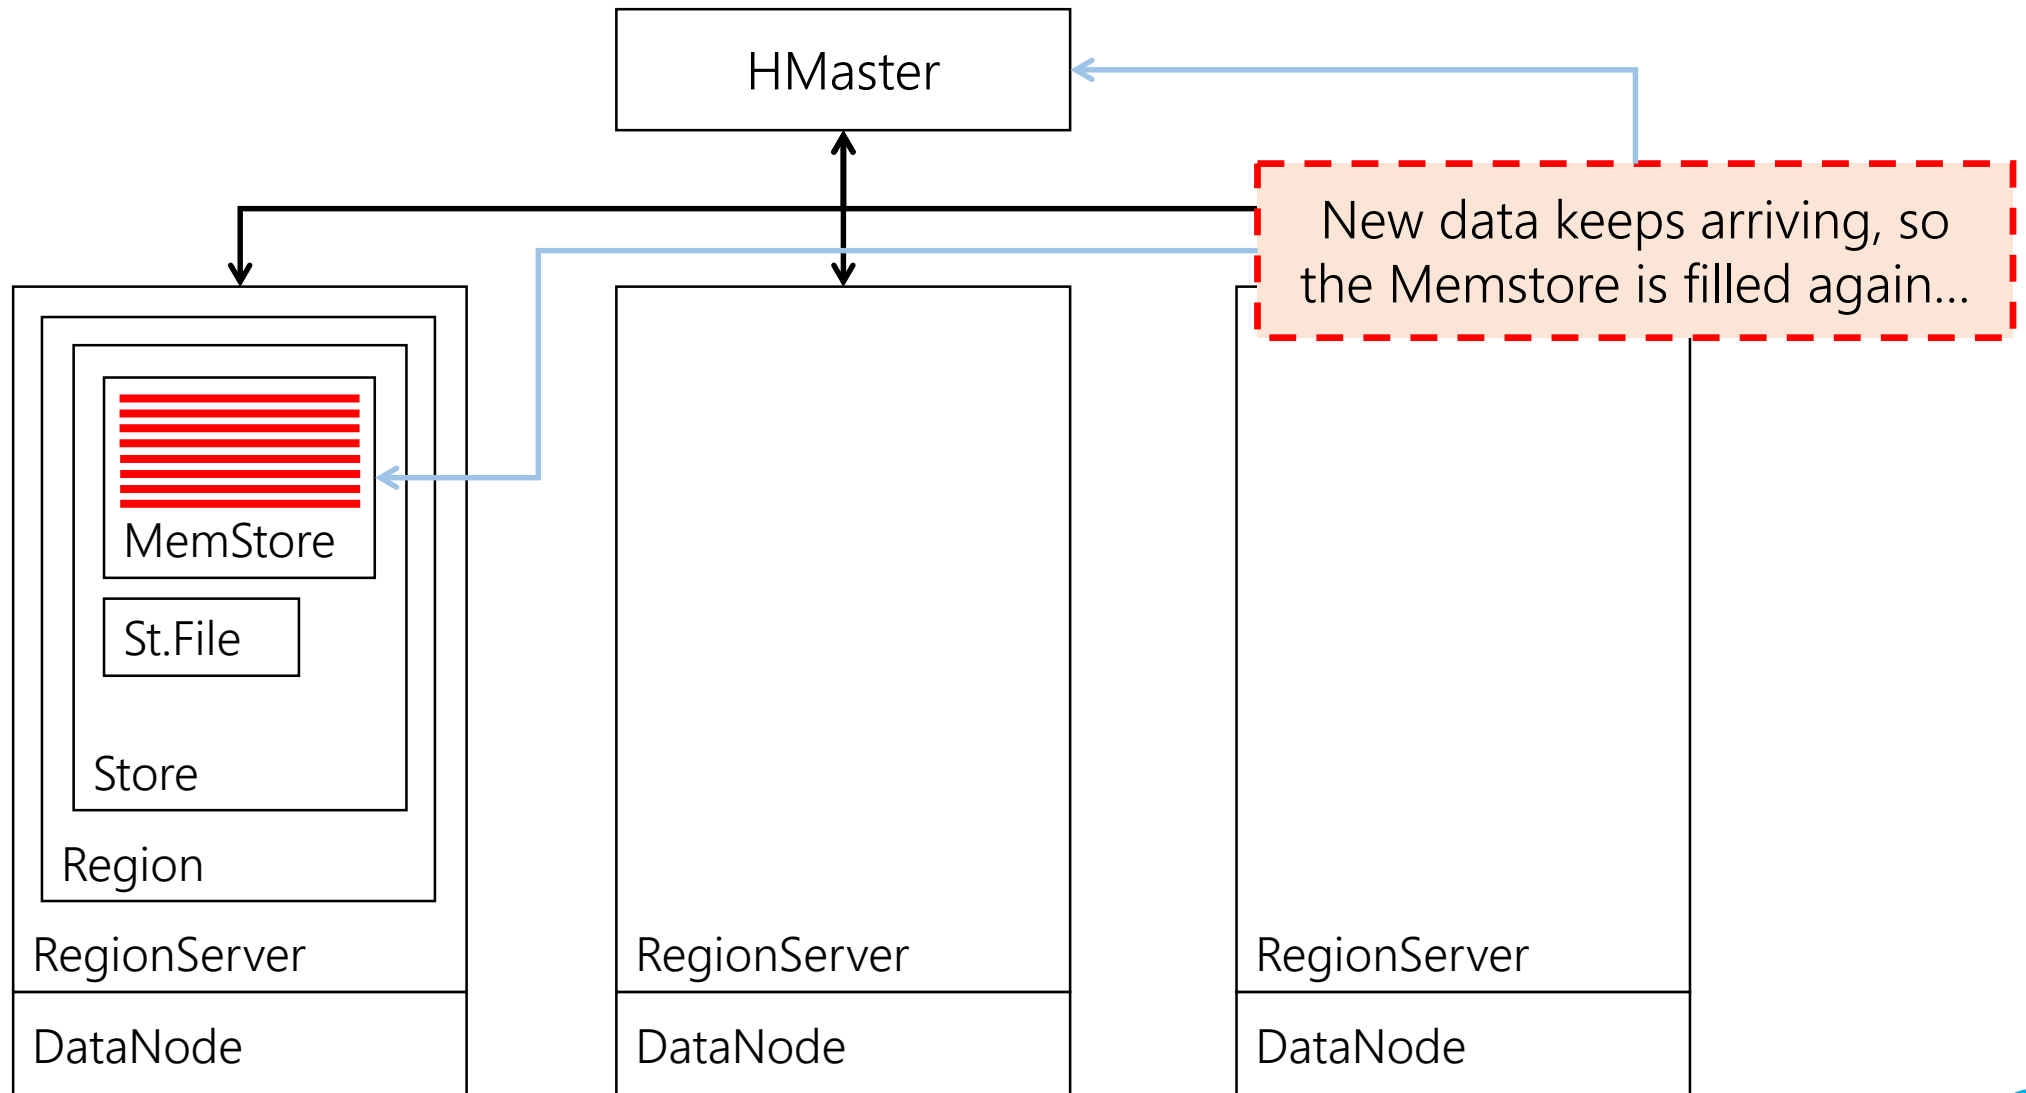
# Example of Flush



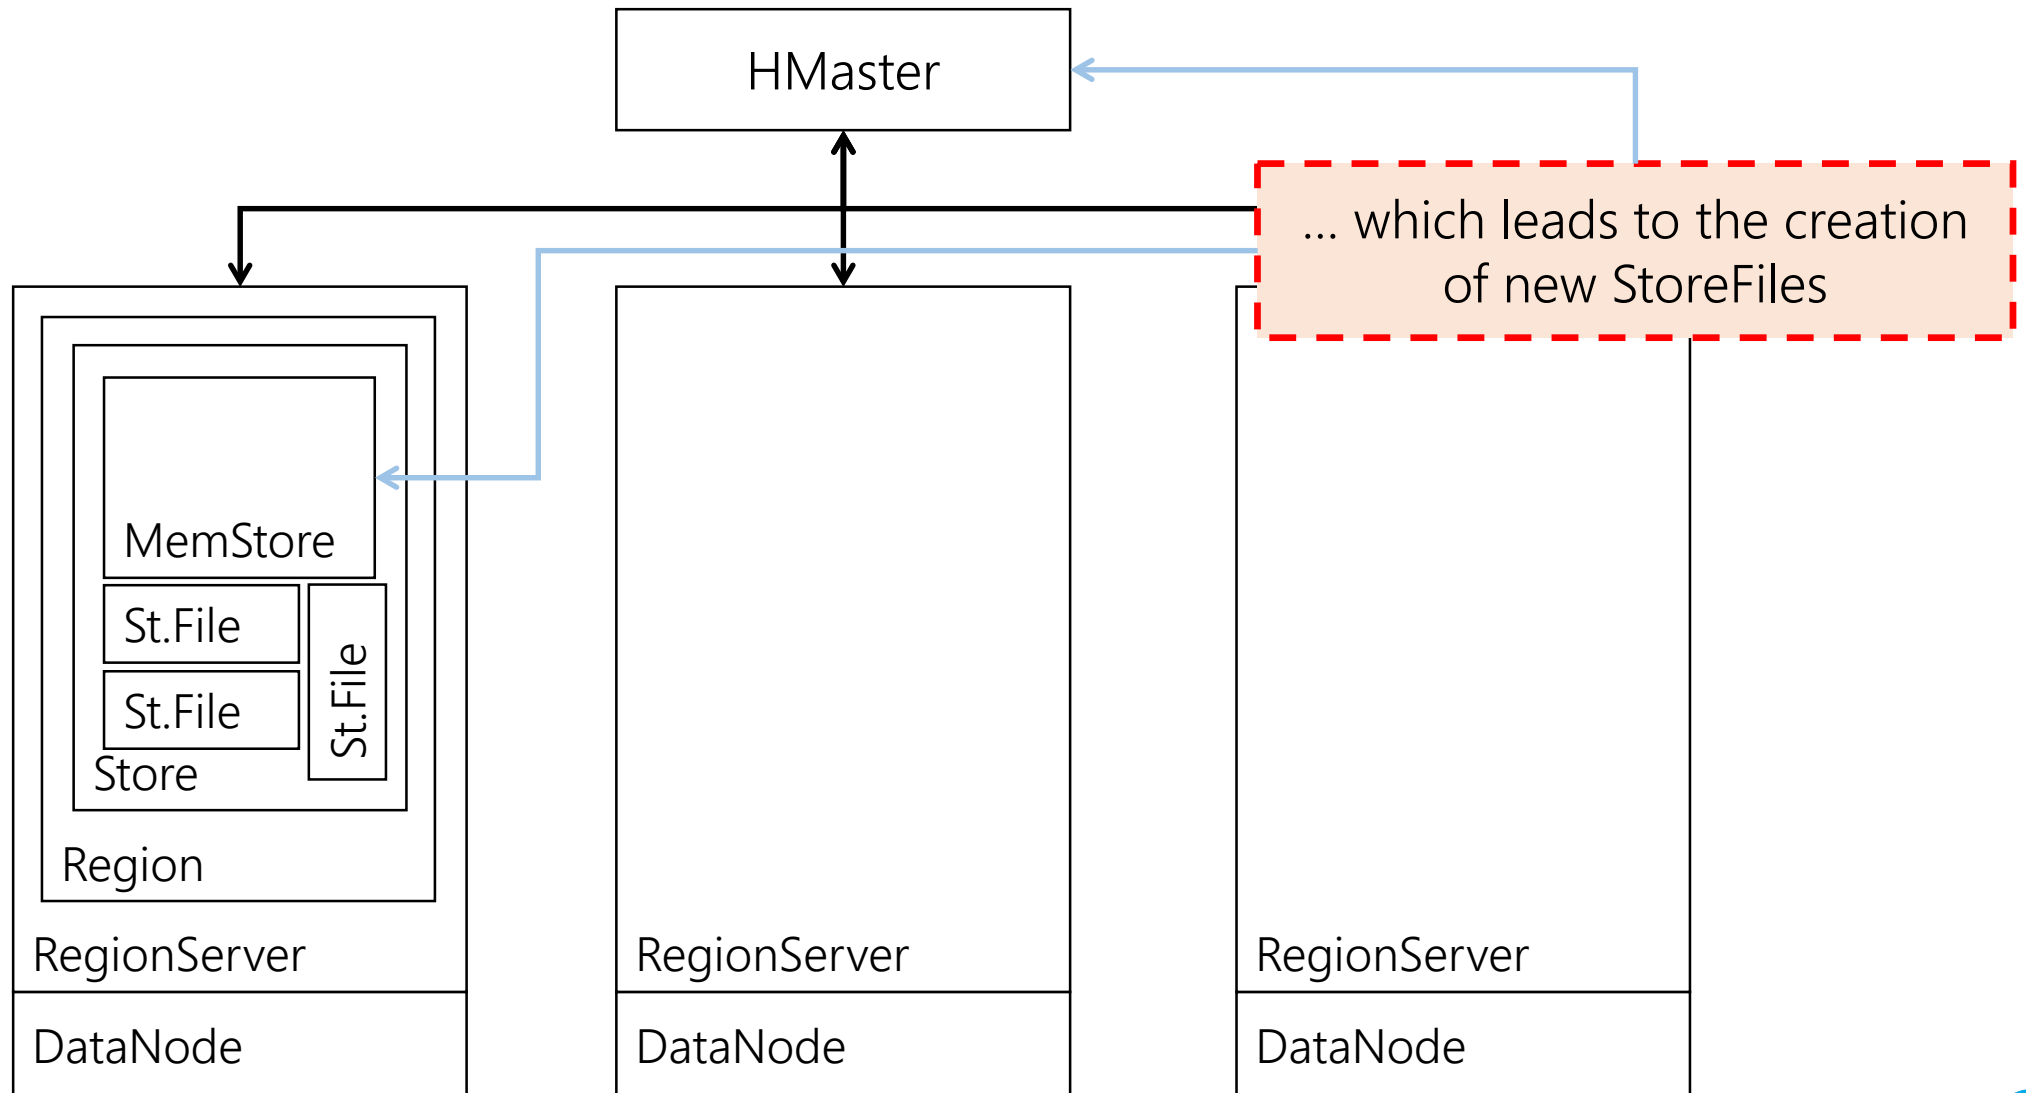
# Example of Flush



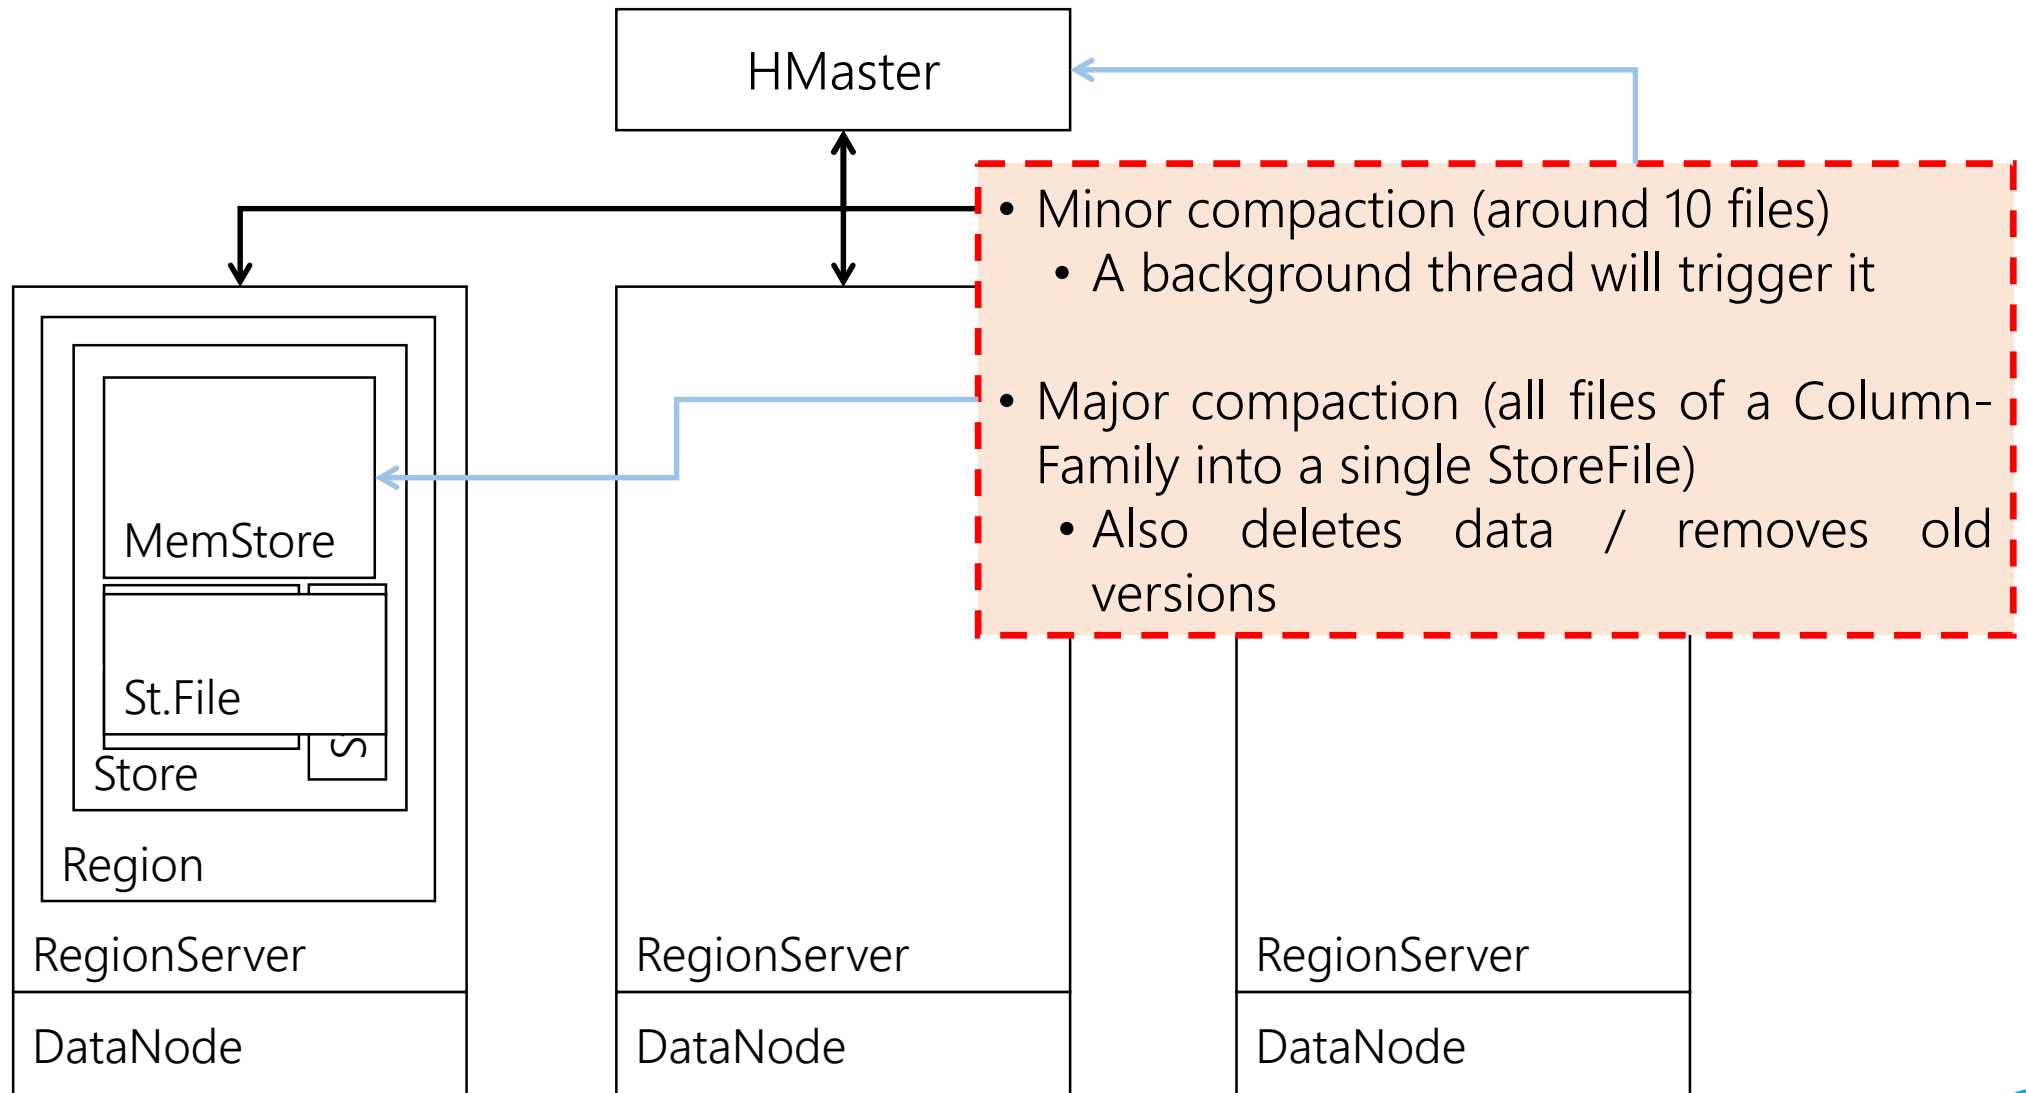
# Example of Compactation



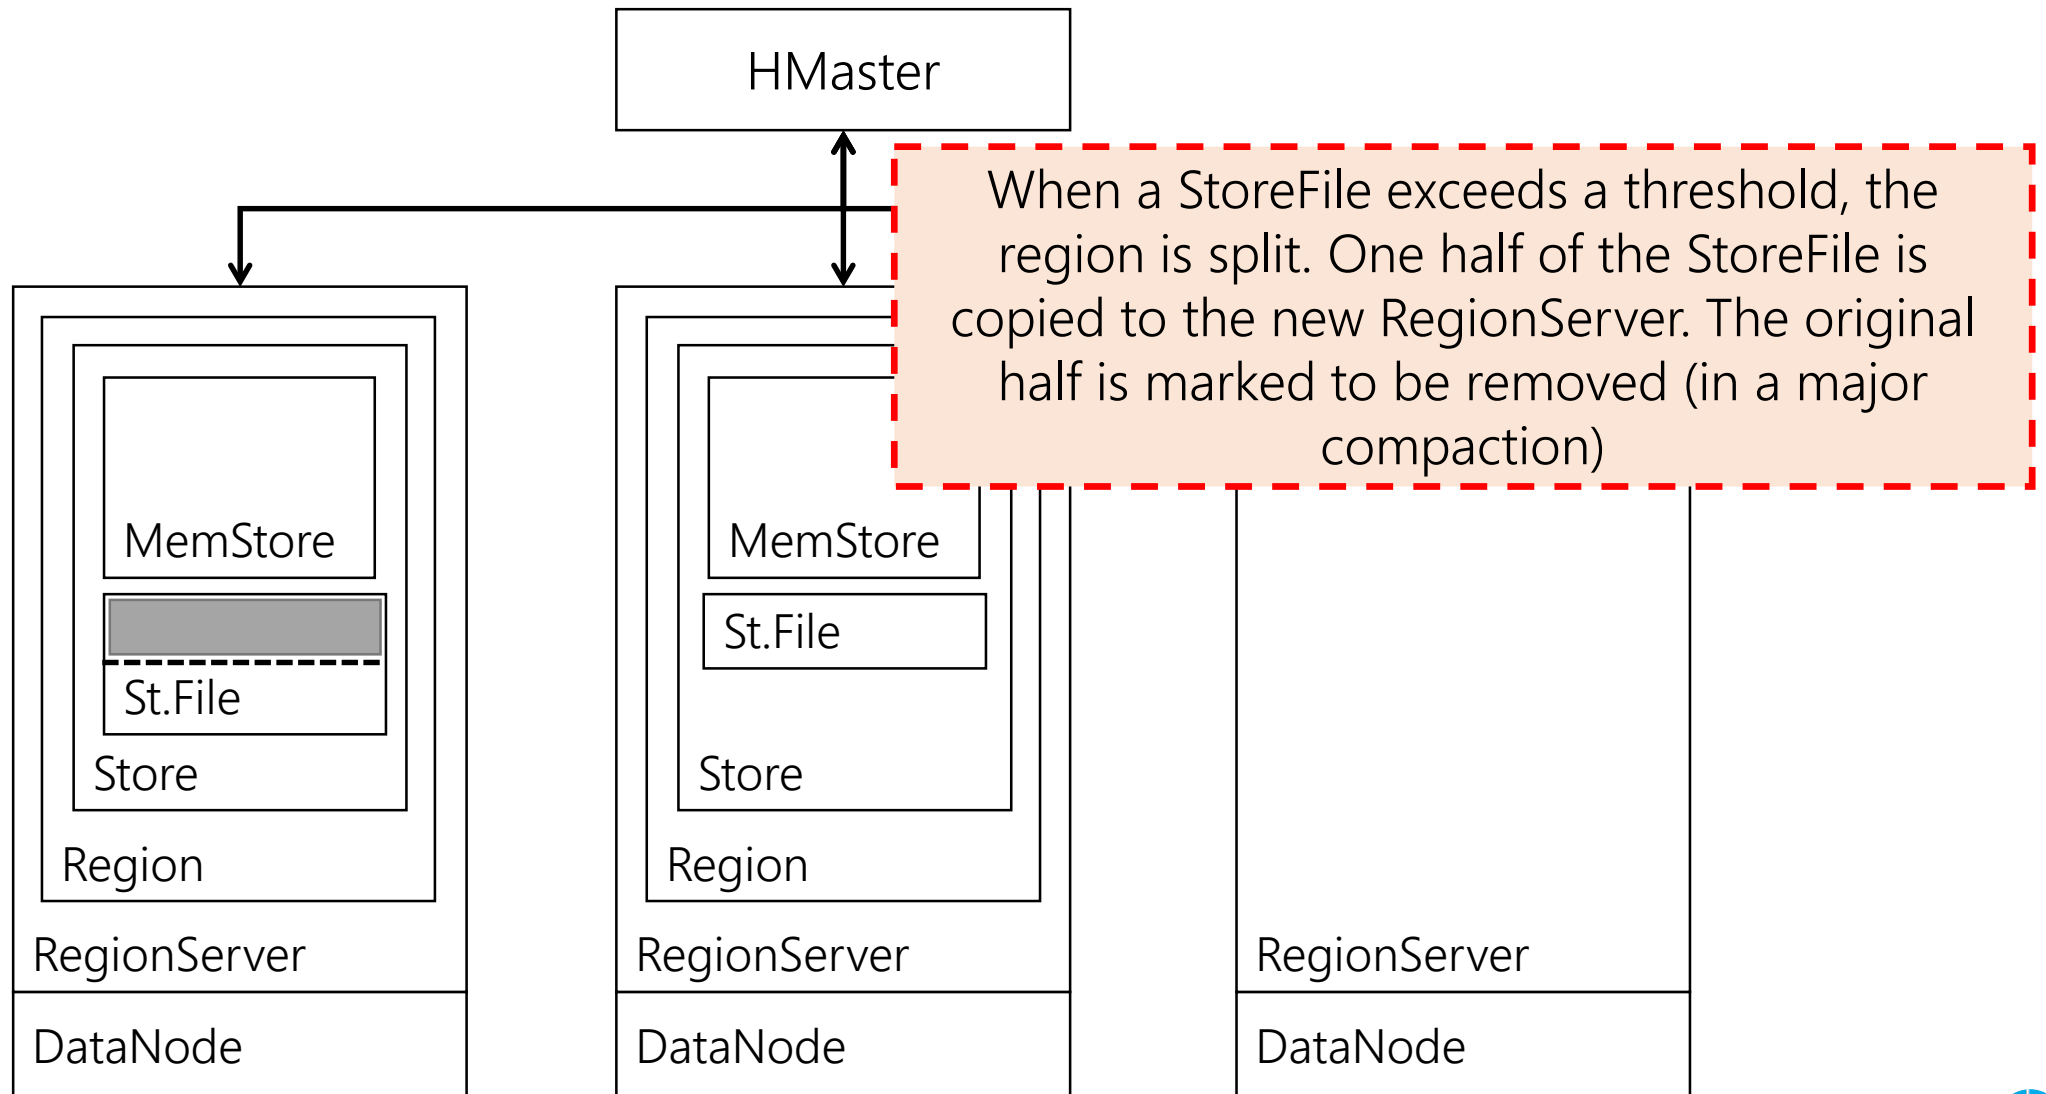
# Example of Compactation



# Example of Compactation



# Example of Split



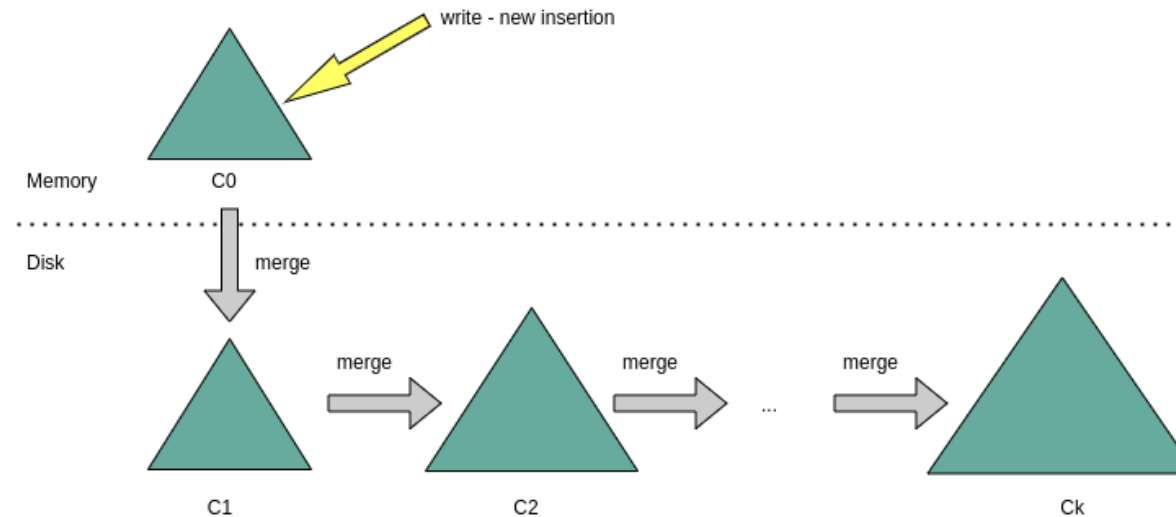
# Underlying Tree Structures

LSM-tree



# Log Structured Merge-trees


- Defers and batches index changes
- Consists on two structures
  - In Memory
    - Not necessarily a B-Tree
    - Node different from disk blocks
  - In Disk
    - Multiple components
      - Blocks 100% full
- Regularly merges trees in one level into the next one



Aina Montalban, based on N. Neeraj

# HBase LSM-tree implementation


- In Memory (MemStore)
  - Holds the most recent updates for fast lookup
  - Sorted by key
- In Disk (StoreFiles)
  - Immutable Sorted-String Tables (with Bloom Filters and Indexes)
    - May contain different versions of the same row
      - All of them need to be accessed at query time
    - Regularly performs a size-tiered merge process
      - Old SSTables not overwritten (available for recovery)



# LSM-tree maintenance algorithms

On memory structure reaching threshold:

1. Take next in memory leafs
2. Flush them to an SSTable (of the established size)



---

On triggering a compactation

1. Take  $n$  SSTables and merge them
2. Put the merge in an in-memory buffer
3. If buffer size exceeds chunk size
  1. Write one chunk to disk
  2. Purge buffer
  3. Keep exceeds in the buffer


# Underlying Hash Structures

Linear hash

Consistent hash

# Distributed Hashing (alternative to a tree)

- Hash does neither support range queries nor nearest neighbours search
- Distributed hashing challenges
  - Dynamicity:
    - Typical hash function
$$h(x) = f(x) \% \text{\#servers}$$
    - Adding a new server implies modifying hash function
      - Massive data transfer
      - Communicating the new function to all servers
  - Location of the hash directory:
    - Any access must go through the hash directory



Aina Montalban, based on S. Abiteboul et al.

“%” is the mathematical modulo operation

# Adding a node in static hash

3 Nodes

Bucket/Node 0

Key 0  
Key 3  
Key 6  
Key 9

Bucket/Node 1

Key 1  
Key 4  
Key 7  
Key 10


Bucket/Node 2

Key 2  
Key 5  
Key 8  
Key 11

Bucket/Node 3

# Distributed Linear Hashing (LH\*)

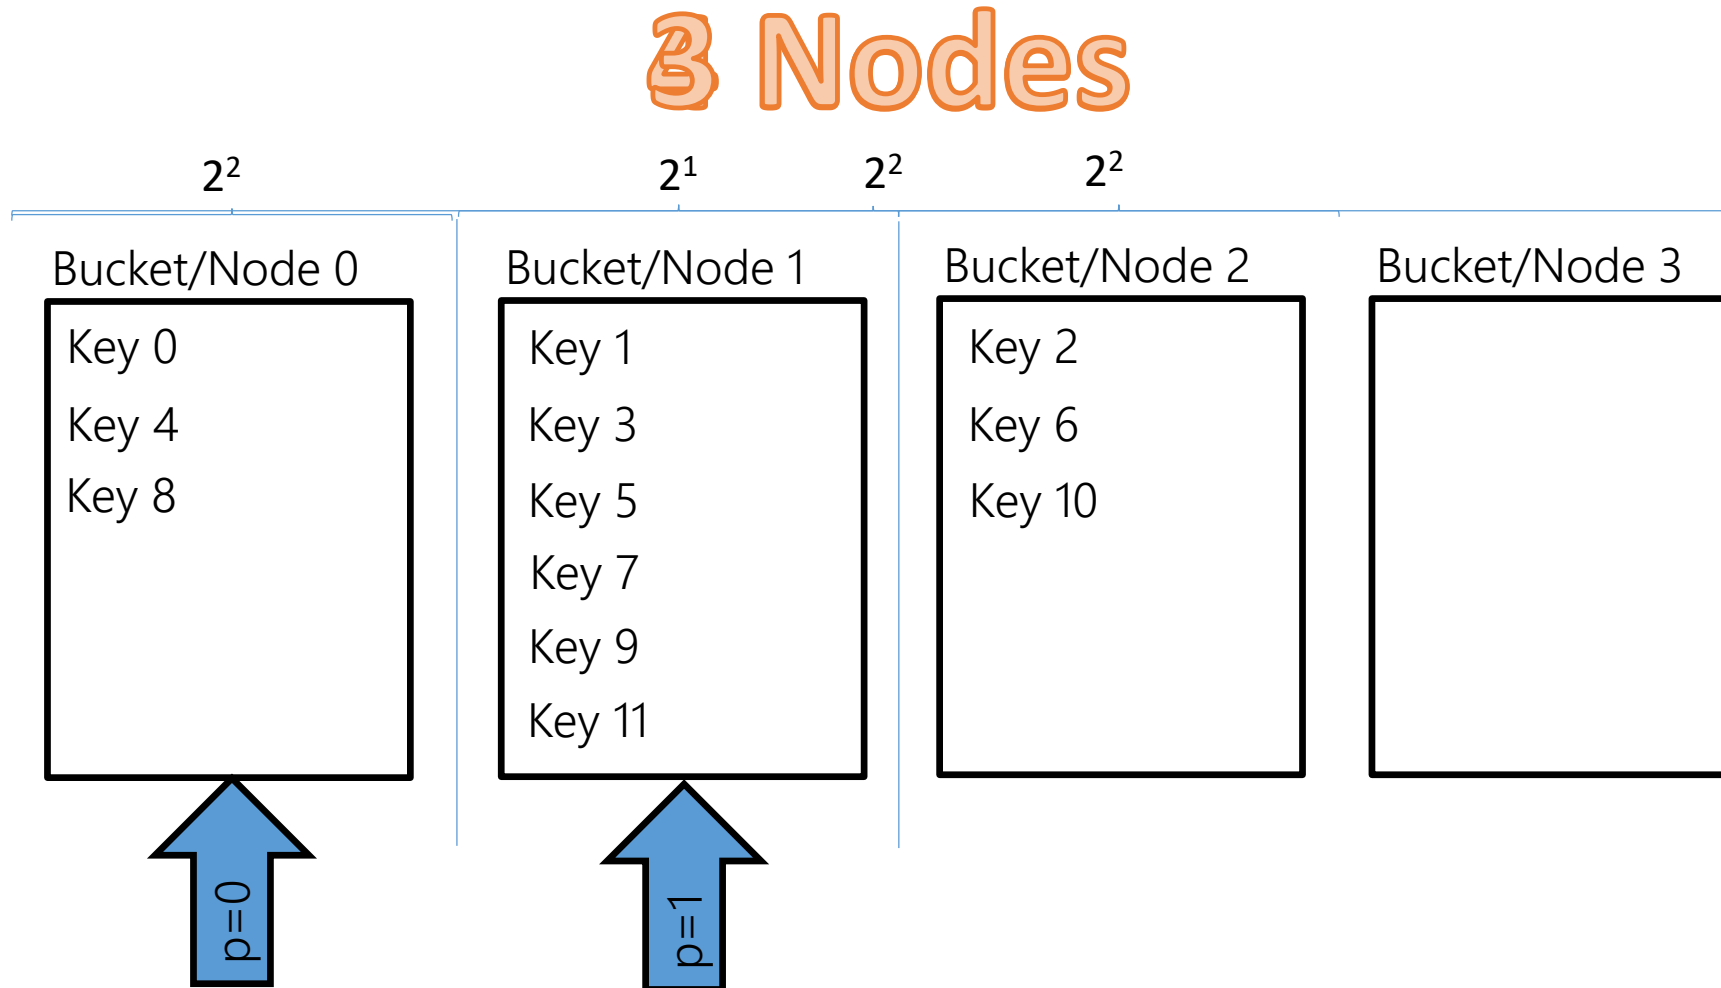
- Maintains an efficient hash in front of dynamicity
  - A split pointer is kept (next bucket to split)
  - A pair of hash functions are considered
    - $\%2^n$  and  $\%2^{n+1}$  (being  $2^n \leq \text{\#servers} < 2^{n+1}$ )
- Overflow buckets are considered
  - When a bucket overflows the pointed bucket splits (not necessarily the overflown one)



Bucket b2 receives a new object

Bucket b0 splits; bucket b2 is linked to a new one


# Adding a node in linear hash






# Consistent Hashing

- The hash function never changes
  - Choose a very large domain  $D$ 
    - Map server IP addresses and object keys to such domain
  - Organize  $D$  as a ring in clockwise order
    - Each node has a successor
  - Objects are assigned as follows:
    - For an object  $o$ ,  $f(o) = D_o$
    - Let  $D_{o'}$  and  $D_{o''}$  be the two nodes in the ring such that
      - $D_{o'} < D_o \leq D_{o''}$
    - $o$  is assigned to  $D_{o'}$
- Further refinements:
  - Assign to the same server several hash values (virtual servers) to balance load
  - Same considerations for the hash directory as for LH\*



# Adding a new server in Consistent Hashing




S. Abiteboul et al.

- Adding a new server is straightforward
  - It is placed in the ring and part of its successor's objects are transferred to it

# Data Design

## Challenge I

# Logical structure

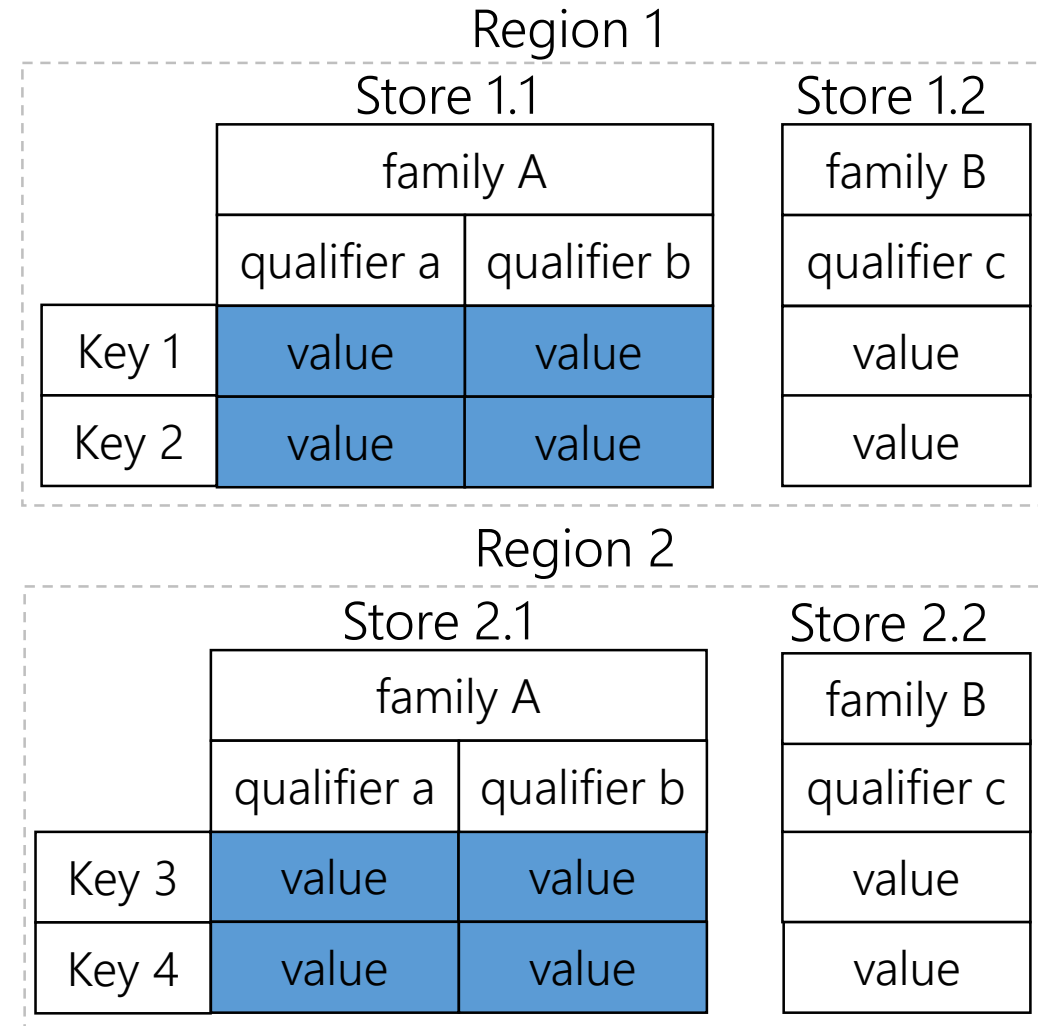



Region 1			
		family A	
		qualifier a	qualifier b
Key 1	value	value	
Key 2	value	value	
		family B	
		qualifier c	
		value	
		value	

Region 2			
		family A	
		qualifier a	qualifier b
Key 3	value	value	
Key 4	value	value	
		family B	
		qualifier c	
		value	
		value	

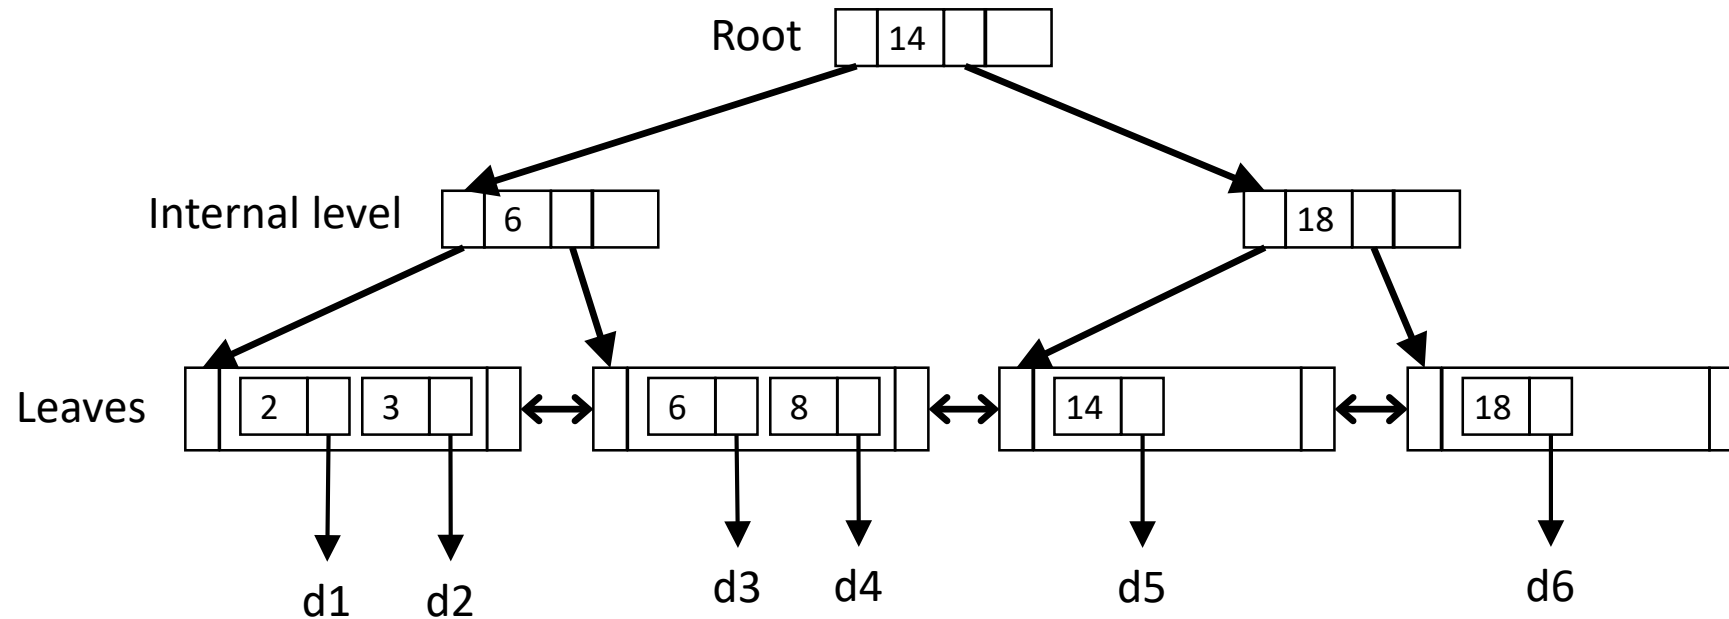
# Physical structure



# Catalog Management

Challenge II

# Metadata hierarchical structure



# HBase Component Roles

- One coordinator server
  - Maintenance of the table schemas
    - Root region
  - Monitoring of services (heartbeating)
  - Assignment of regions to servers
- Many region servers
  - Each handling around 100-1.000 regions
  - Apply concurrency and recovery techniques
  - Managing split of regions
    - Regions can be sent to another server (load balancer)
  - Managing merge of regions
- Client nodes
  - Cache the metadata sent by the region servers



# Metadata hierarchical structure


## Root table

Store locations of metadata tables




## Metadata tables

Store locations of user data tables



...



## User data tables

Store data


key	columns
k1	Row 1
k2	Row 2
...	...
km	Row m

Table T

key	columns
	Row 1
	Row 2
...	...
k'p	Row p

Table T'

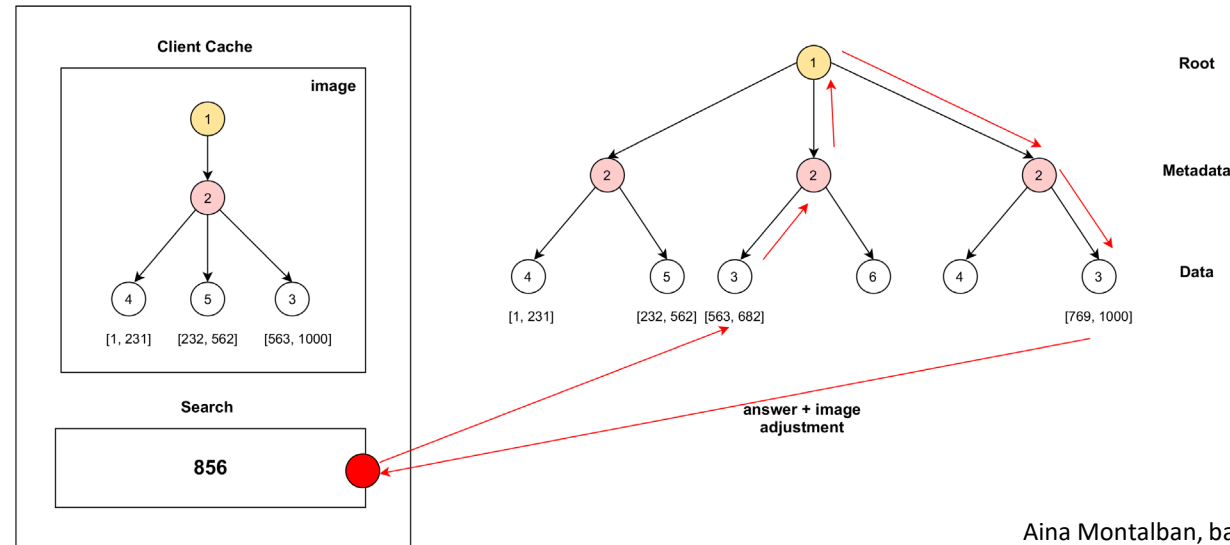
...



more tables

# Metadata synchronization in HBase

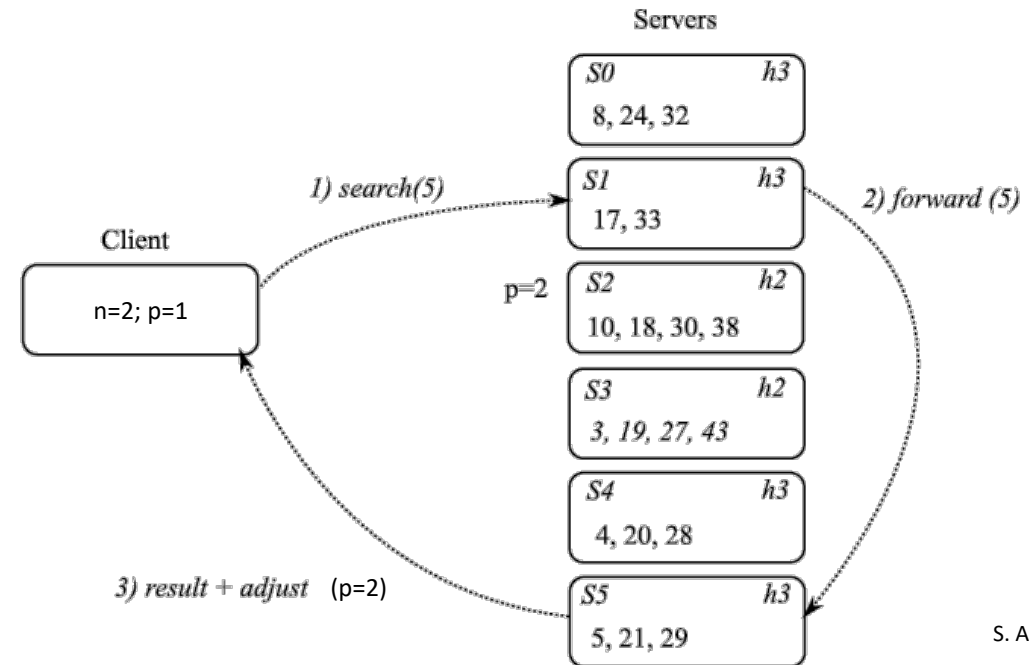
- Split and merge invalidate the cached metadata
  - a) Gossiping
  - b) Lazy updates
- Discrepancies may cause out-of-range errors, which trigger a stabilization protocol (i.e., **mistake compensation**)
  - Apply forwarding path
    - If an out-of-range error is triggered, it is forwarded upward
    - In the worst case (i.e., reaching the root), 6 network round trips



Aina Montalban, based on S. Abiteboul et al.

# Updating the Hash Directory in LH\*

- Traditionally, each participant has a copy of the hash directory
  - Changes in the hash directory (either hash functions or splits) imply *gossiping*
    - Including clients nodes
    - It might be acceptable if not too dynamic
- Alternatively, they may contain a partial representation and assume lazy adjustment
  - Apply forwarding path



# Transaction Management

Challenge III

# Single row guarantees


- Atomicity
  - Only single row guarantees (even across families)
    - ... since different families of a row are not distributed
- Consistency
  - Replication relies on HDFS
- Isolation
  - Locking mechanism at row level
    - Does not guarantee snapshot isolation (only read committed)
      - Rows (all families) are separately consistent
      - Sets of rows may not be consistent as a whole
- Durability
  - Write Ahead Log implementation

<https://hbase.apache.org/acid-semantics.html>


# Query processing

Challenge IV

# Global execution



# Local execution





# Closing

# Summary

- Key-value abstraction
- HBase functional components
- Directory caching
  - Mistake compensation
- Data distribution structures
  - LSM-Tree
  - Linear hash
  - Consistent hash

# References

- P. O'Neil et al. *The log-structured merge-tree (LSM-tree)*. Acta Informatica, 33(4). Springer, 1996
- F. Chang et al. *Bigtable: A Distributed Storage System for Structured Data*. OSDI'06
- S. Abiteboul et al. *Web Data Management*. Cambridge University Press, 2011. <http://webdam.inria.fr/Jorge>
- N. Neeraj. *Mastering Apache Cassandra*. Packt, 2015
- O. Romero et al. *Tuning small analytics on Big Data: Data partitioning and secondary indexes in the Hadoop ecosystem*. Information Systems, 54. Elsevier, 2016
- A. Petrov. *Algorithms Behind Modern Storage Systems*. Communications of the ACM 61(8), 2018